

Evolutionary database design: Refactoring databases



Pramod Sadalage
@pramodsadalage
ThoughtWorks Inc.

SSID: Pramods Wifi Network
Password: Database#1

Motivations

- Continuous Delivery
- Collaboration - data team
- Changing requirements
- Devops - data team
- Deploying database changes

Why Continuous Delivery?

To get fast feedback from
users, release frequently

Reduce risk of releasing

Learning and responding to
customer needs is critical

Achieve Continuous Delivery

Close collaboration between
everyone involved in delivery

Extensive automation of all
parts involved in delivery
process

Reduce the gap between
development and operations






















CD should involve the
database team

**How does continuous
delivery apply to
databases?**

Collaboration techniques

Pair the DBA's and
Developers

Ve

Name	
	profilesrvr.bat
	build.properties.sample
	Rakefile
	report.cmd
▶ 	classes
▶ 	config
▼ 	db
▶ 	admin
▶ 	backup
▶ 	DataModel
▶ 	deltas
▶ 	lib
▶ 	migration
▶ 	PerformanceDeltas
▶ 	QA-tabledata-notready
▶ 	QADeltas
▶ 	QATools
▶ 	refactoring
▶ 	tools
	cleandb.sql
	comments.sql

se

Lets see our code

`github.com/pramodsadalage/database-refactoring`

Your Projects

- Is your project a single database project
- How many applications talk to your database?
- Does your app depend on multiple databases?
- Database dependency in version control?

Automation techniques

Let developers provision
application database

```
cd-databases:>git clone https://github.com/pramodsadalage/evodb.git
Cloning into 'evodb'...
remote: Counting objects: 115, done.
remote: Total 115 (delta 0), reused 0 (delta 0)
Receiving objects: 100% (115/115), 7.59 MiB | 1.73 MiB/s, done.
Resolving deltas: 100% (14/14), done.
cd-databases:>cd evodb/
evodb:>cp build.properties.template build.properties
evodb:>vi build.properties
evodb:>ant create_schema
Buildfile: /Users/Thoughtworker/cd-databases/evodb/build.xml
```

```
create_schema:
    [echo] Admin UserName: system
    [echo] Creating Schema: malmo
    [sql] Executing commands
    [sql] 4 of 4 SQL statements executed successfully
```

```
BUILD SUCCESSFUL
Total time: 0 seconds
evodb:>
```

```
evodb:>ant dbinit
Buildfile: /Users/Thoughtworker/cd-databases/evodb/build.xml
```

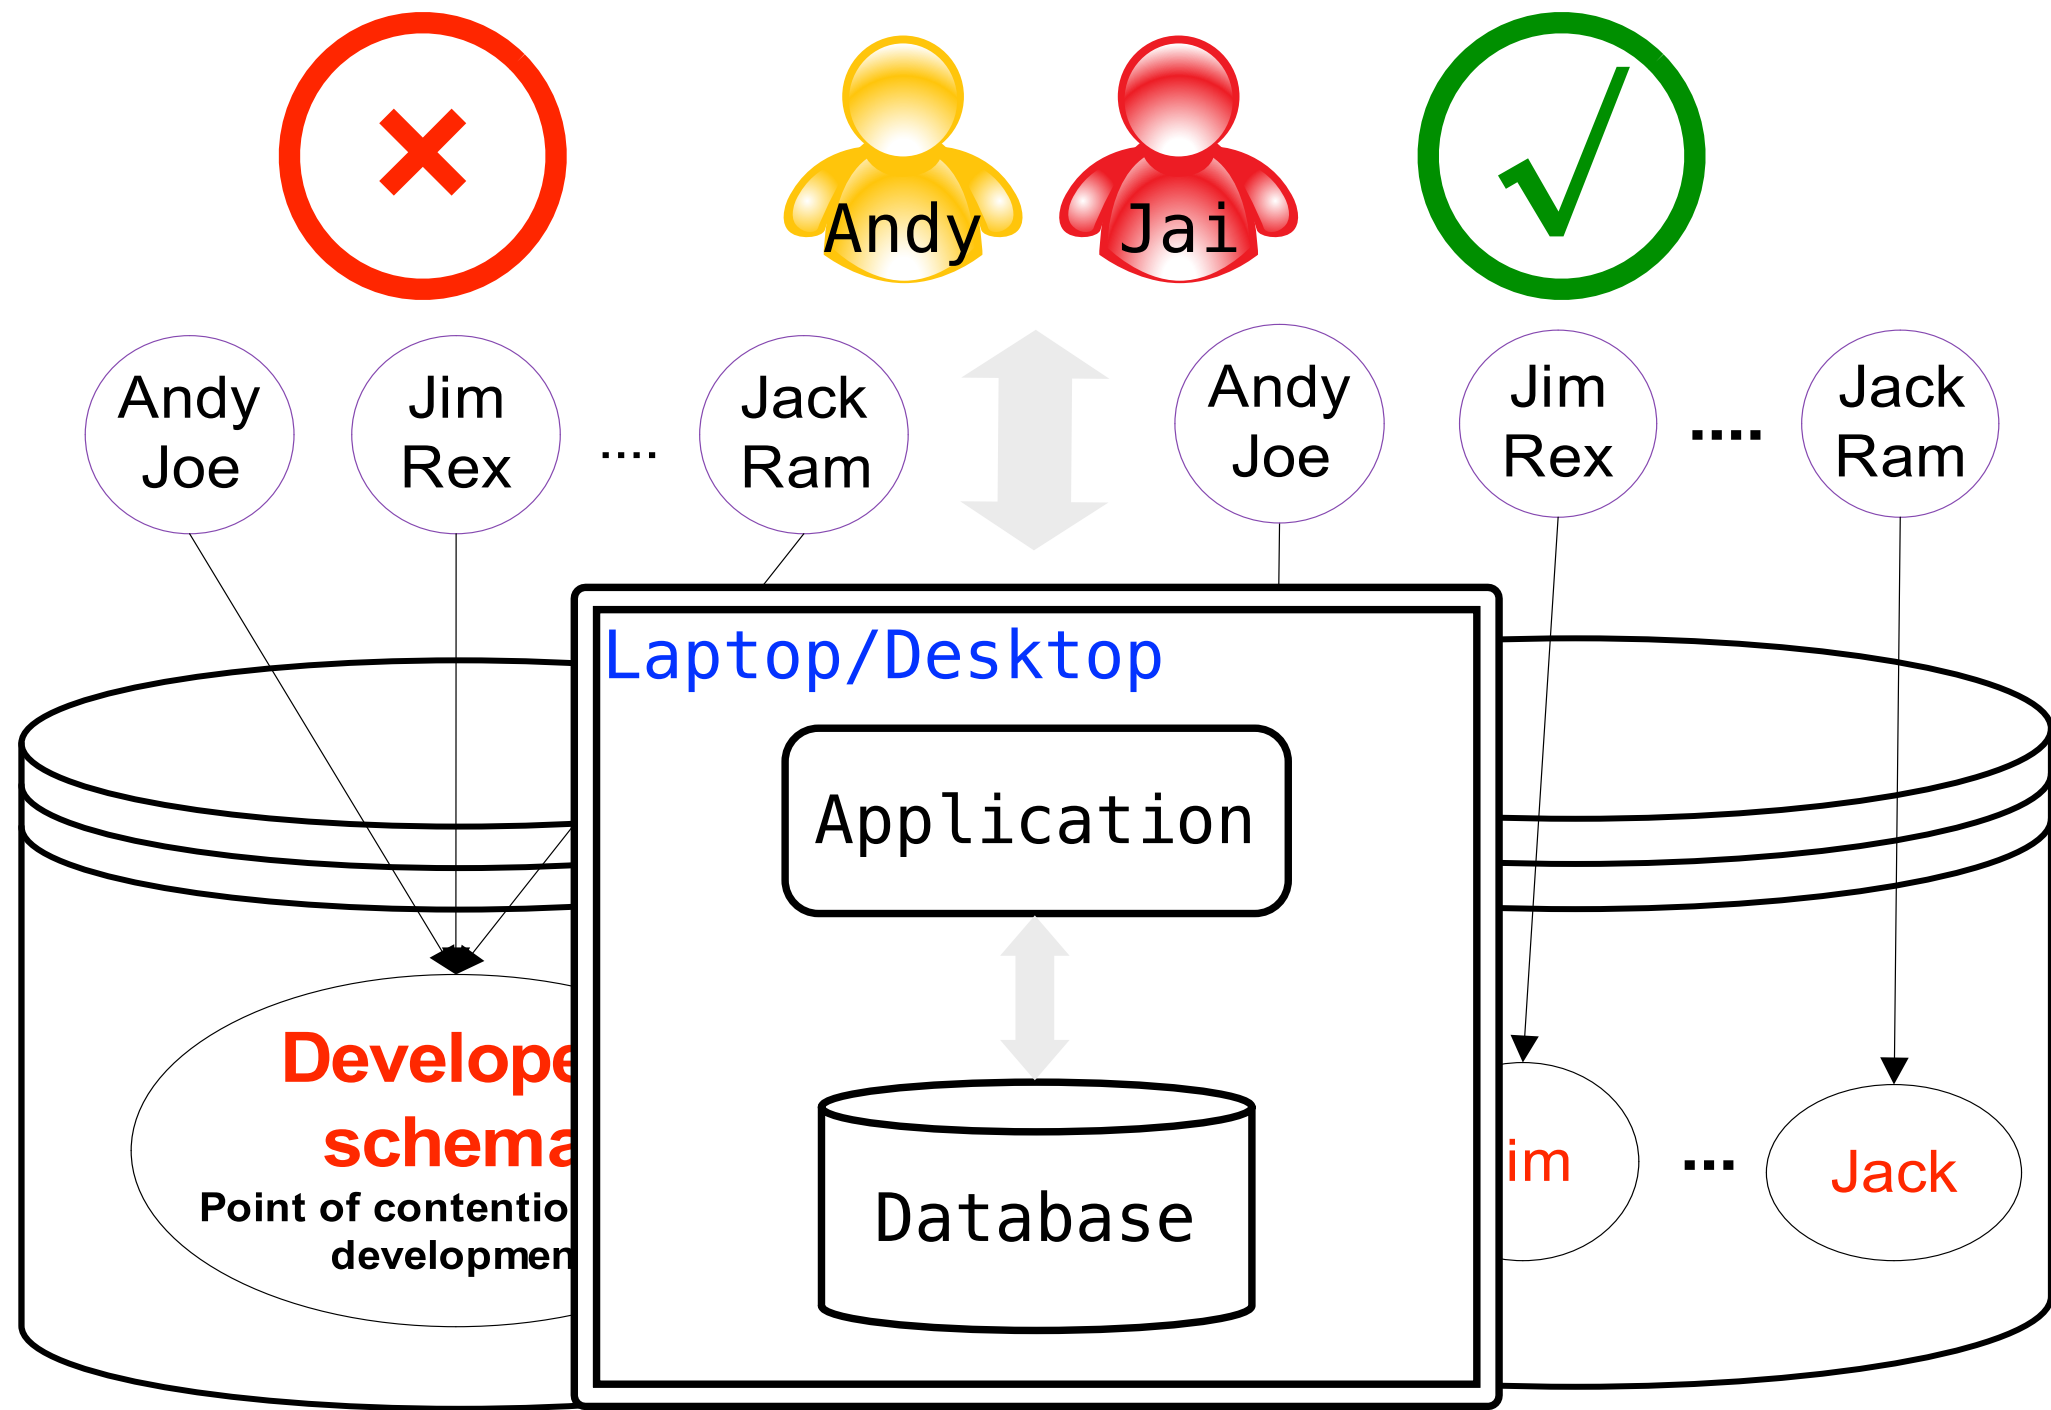
```
dbinit:
  [echo] Working UserName: malmo
```

```
upgrade:
  [dbdeploy] dbdeploy 3.0M3
  [dbdeploy] Reading change scripts from directory /Users/Thoughtworker/cd-databases/evodb/db/migration...
  [dbdeploy] Changes currently applied to database:
  [dbdeploy]   (none)
  [dbdeploy] Scripts available:
  [dbdeploy]   1..7
  [dbdeploy] To be applied:
  [dbdeploy]   1..7
  [dbdeploy] Applying #1: 001IntroduceEmployee.sql...
  [dbdeploy]   -> statement 1 of 2...
  [dbdeploy]   -> statement 2 of 2...
  [dbdeploy] Applying #2: 002IntroduceCustomer.sql...
  [dbdeploy] Applying #3: 003IntroduceAccount.sql...
  [dbdeploy]   -> statement 1 of 2...
  [dbdeploy]   -> statement 2 of 2...
  [dbdeploy] Applying #4: 004CreateCustomerOpenAccountViewAndIndexes.sql...
  [dbdeploy] Applying #5: 005CreateCommentsOnEmployeeAccountCustomerAndContactTable.sql...
  [dbdeploy] Applying #6: 006IntroduceAccountTypeCustomerTypeData.sql...
  [dbdeploy] Applying #7: 007IntroduceAccountTxnAndCustomerTaxID.sql...
  [dbdeploy]   -> statement 1 of 3...
  [dbdeploy]   -> statement 2 of 3...
  [dbdeploy]   -> statement 3 of 3...
```

```
create_test_data:
  [sql] Executing resource: /Users/Thoughtworker/cd-databases/evodb/db/testdata/customer.sql
  [sql] Executing resource: /Users/Thoughtworker/cd-databases/evodb/db/testdata/employee.sql
  [sql] 2 of 2 SQL statements executed successfully
```

```
create_db_code:
  [sql] Executing resource: /Users/Thoughtworker/cd-databases/evodb/db/code/triggers.sql
  [sql] 0 of 0 SQL statements executed successfully
```

```
BUILD SUCCESSFUL
Total time: 1 second
evodb:>
```



Lets build our schema

```
copy template.build.properties to  
build.properties
```

```
Edit build.properties
```

```
ant createschema
```

```
ant dropschema
```


Lets populate our schema

```
./flyway migrate
```

```
./flyway info
```

```
./flyway validate
```

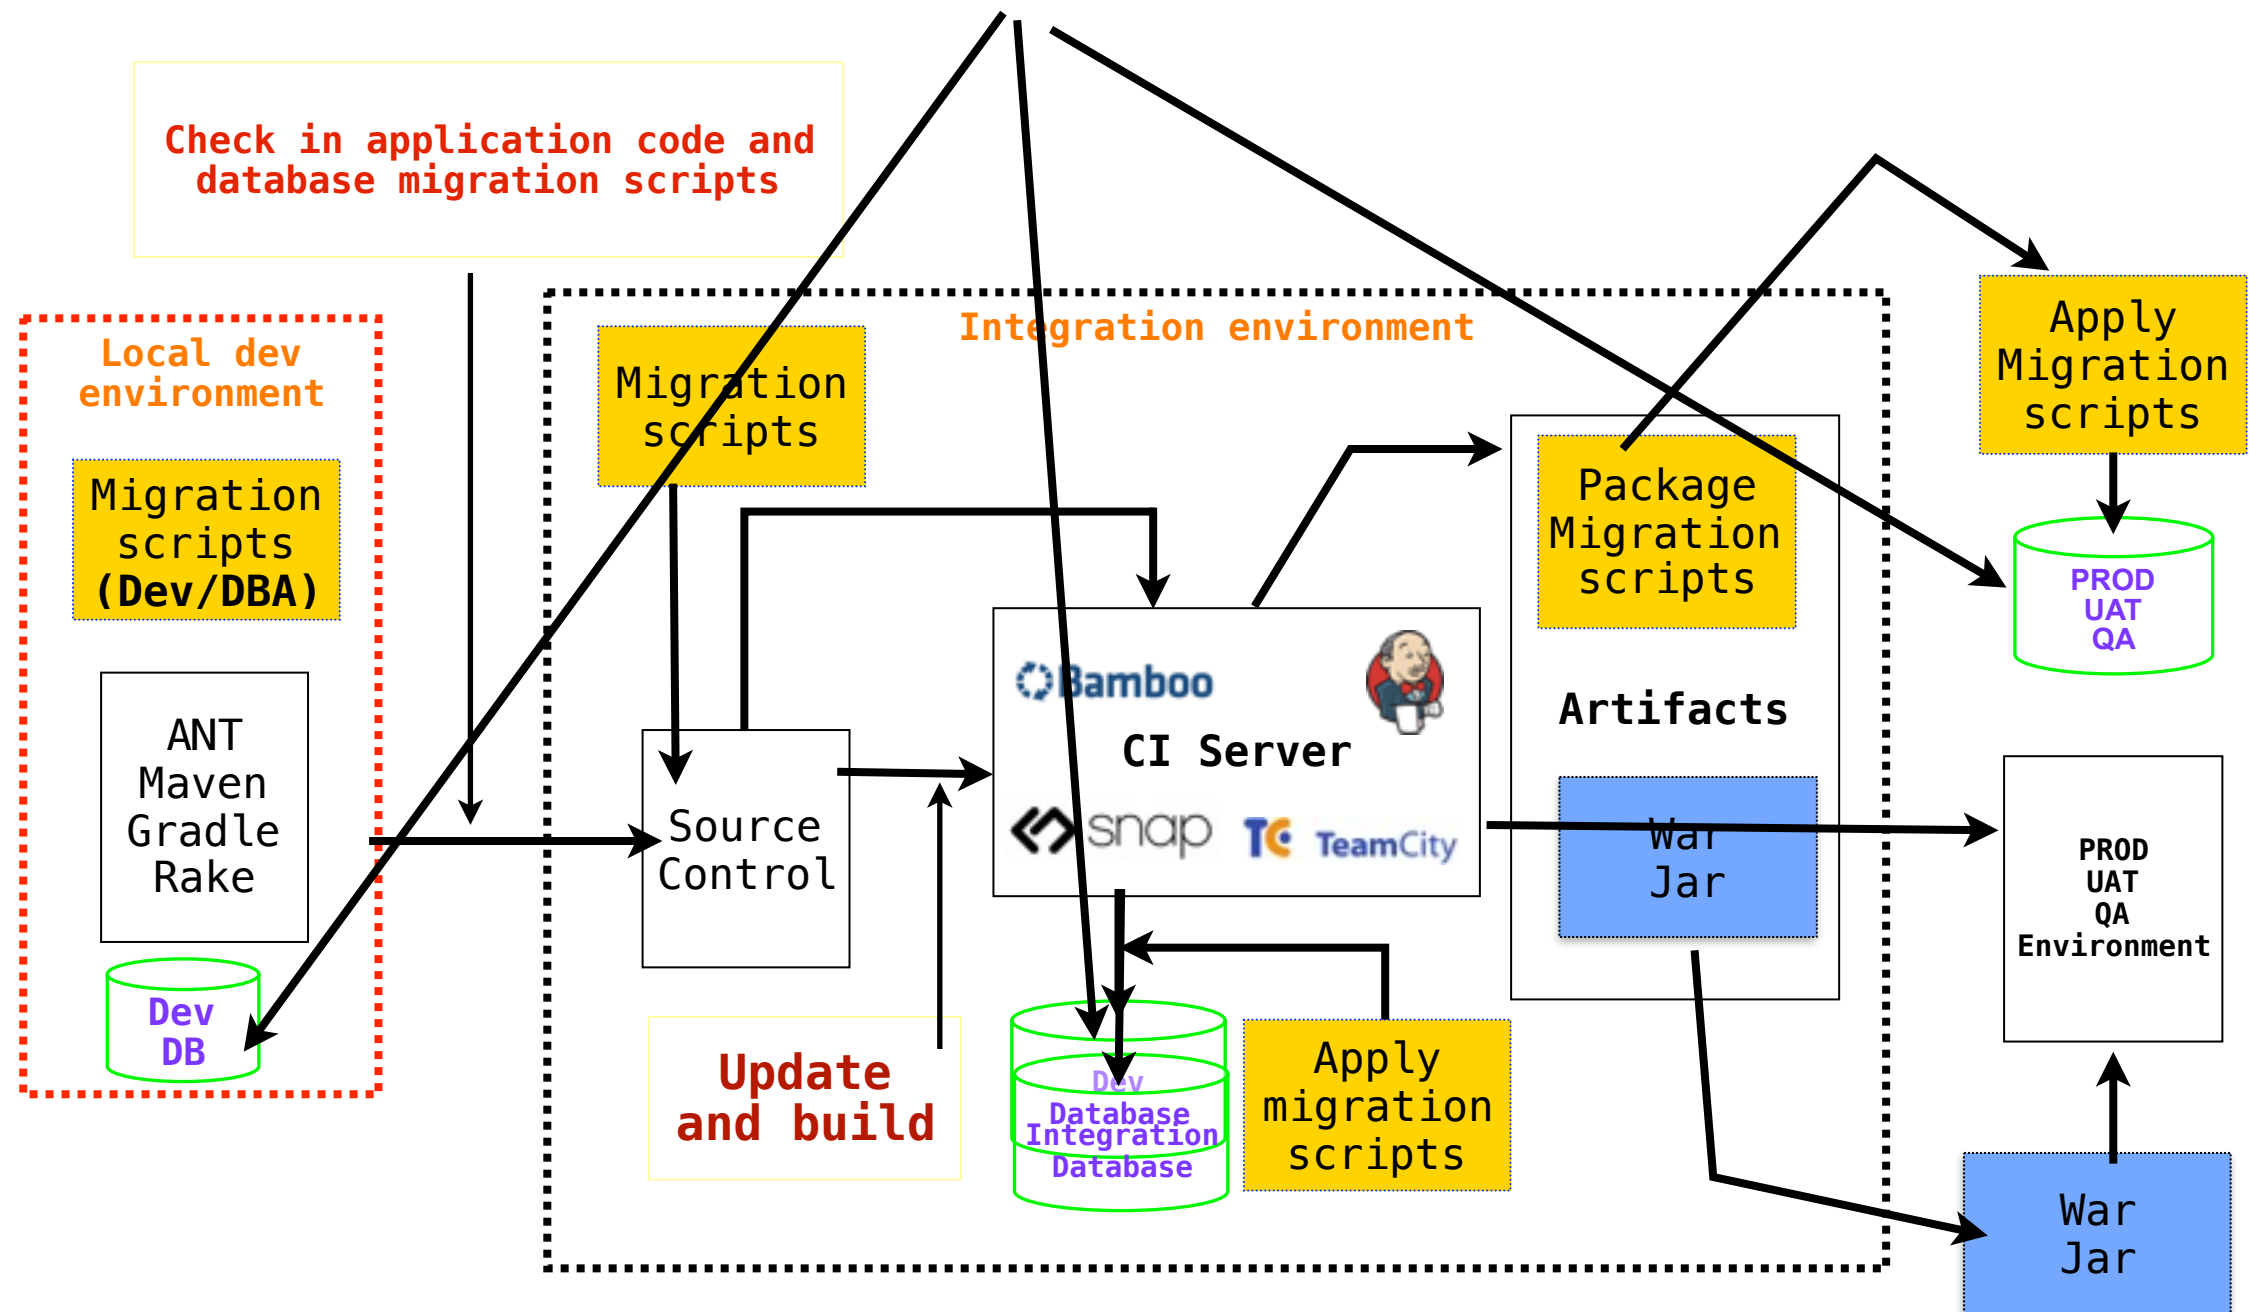
```
ant dbset? what does it do
```

Existing Legacy Database

- Extract the schema and make that your baseline
- How complex is your setup?
- Too many databases, schemas, linked schemas etc?
- Is the physical architecture too complex?

Continuously Integrate
database changes

Database changes applied by DBA Team



Benefits of CI with databases

- Test application code and database at one place
- Generate code and database artifacts
- Integrate application and database changes in an independent environment
- Show current state of application and database to all

DB artifacts in CI

Project EvoDB



Workspace



Last Successful Artifacts



application.jar

41.48 KB



dbupgrade.zip

3.01 MB



Recent Changes



Latest Test Result (no failures)

application artifact for
build <<n>>

database artifact for
build <<n>>

Tracking Changes

- Each change is a delta/migration script
- Migration scripts are development time activity not deployment time tasks
- Package migration scripts for automated deployment
- Same scripts for: developers, QA, UAT and Production

http://10.60.0.75:8080/

Lets one of us make a change and push.

Apply change locally

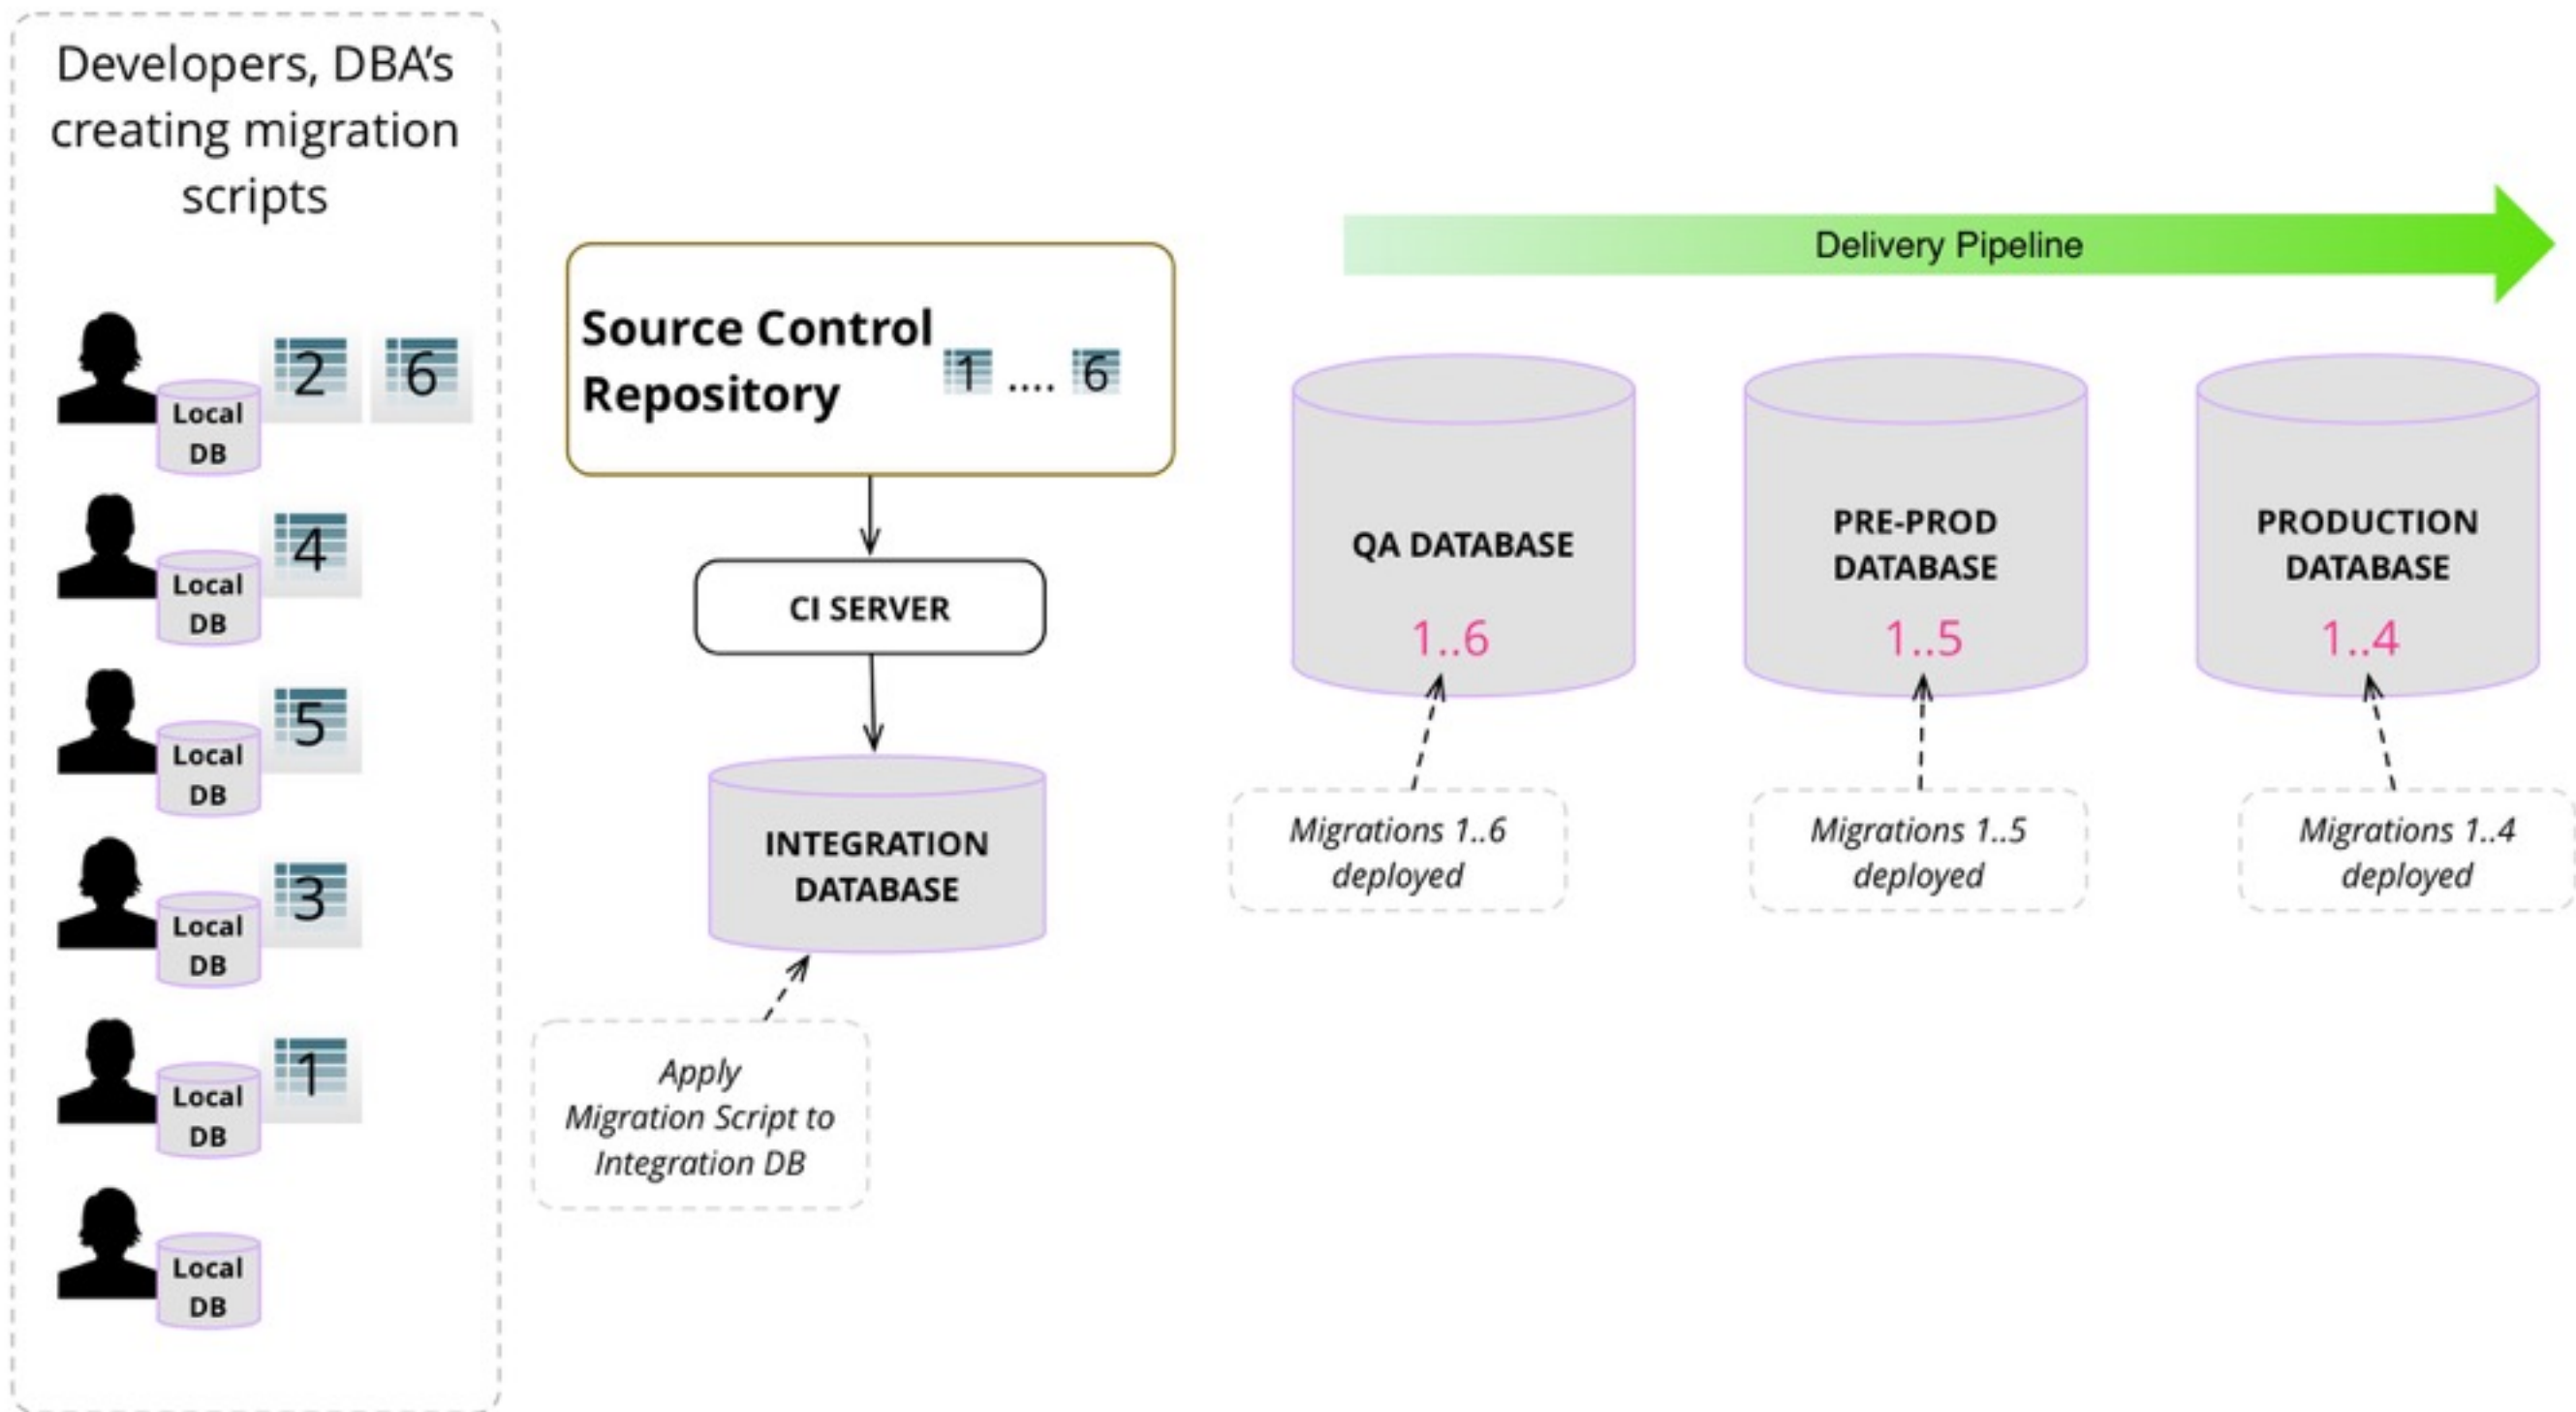
See the rest of the app works with change

Commit and Push.

See Jenkins work

How to deal with conflicts, multiple projects using same database, large teams

Tracking Changes



Deployment

- Database migration/upgrade should be a development time task not deployment time task
- Package all the migration scripts, during Continuous Integration cycle
- Apply the migration scripts
- Deploy frequently to reduce risk

Shifting to code

Download artifacts from Jenkins

What does it have?

Can you apply all the changes to a different schema?

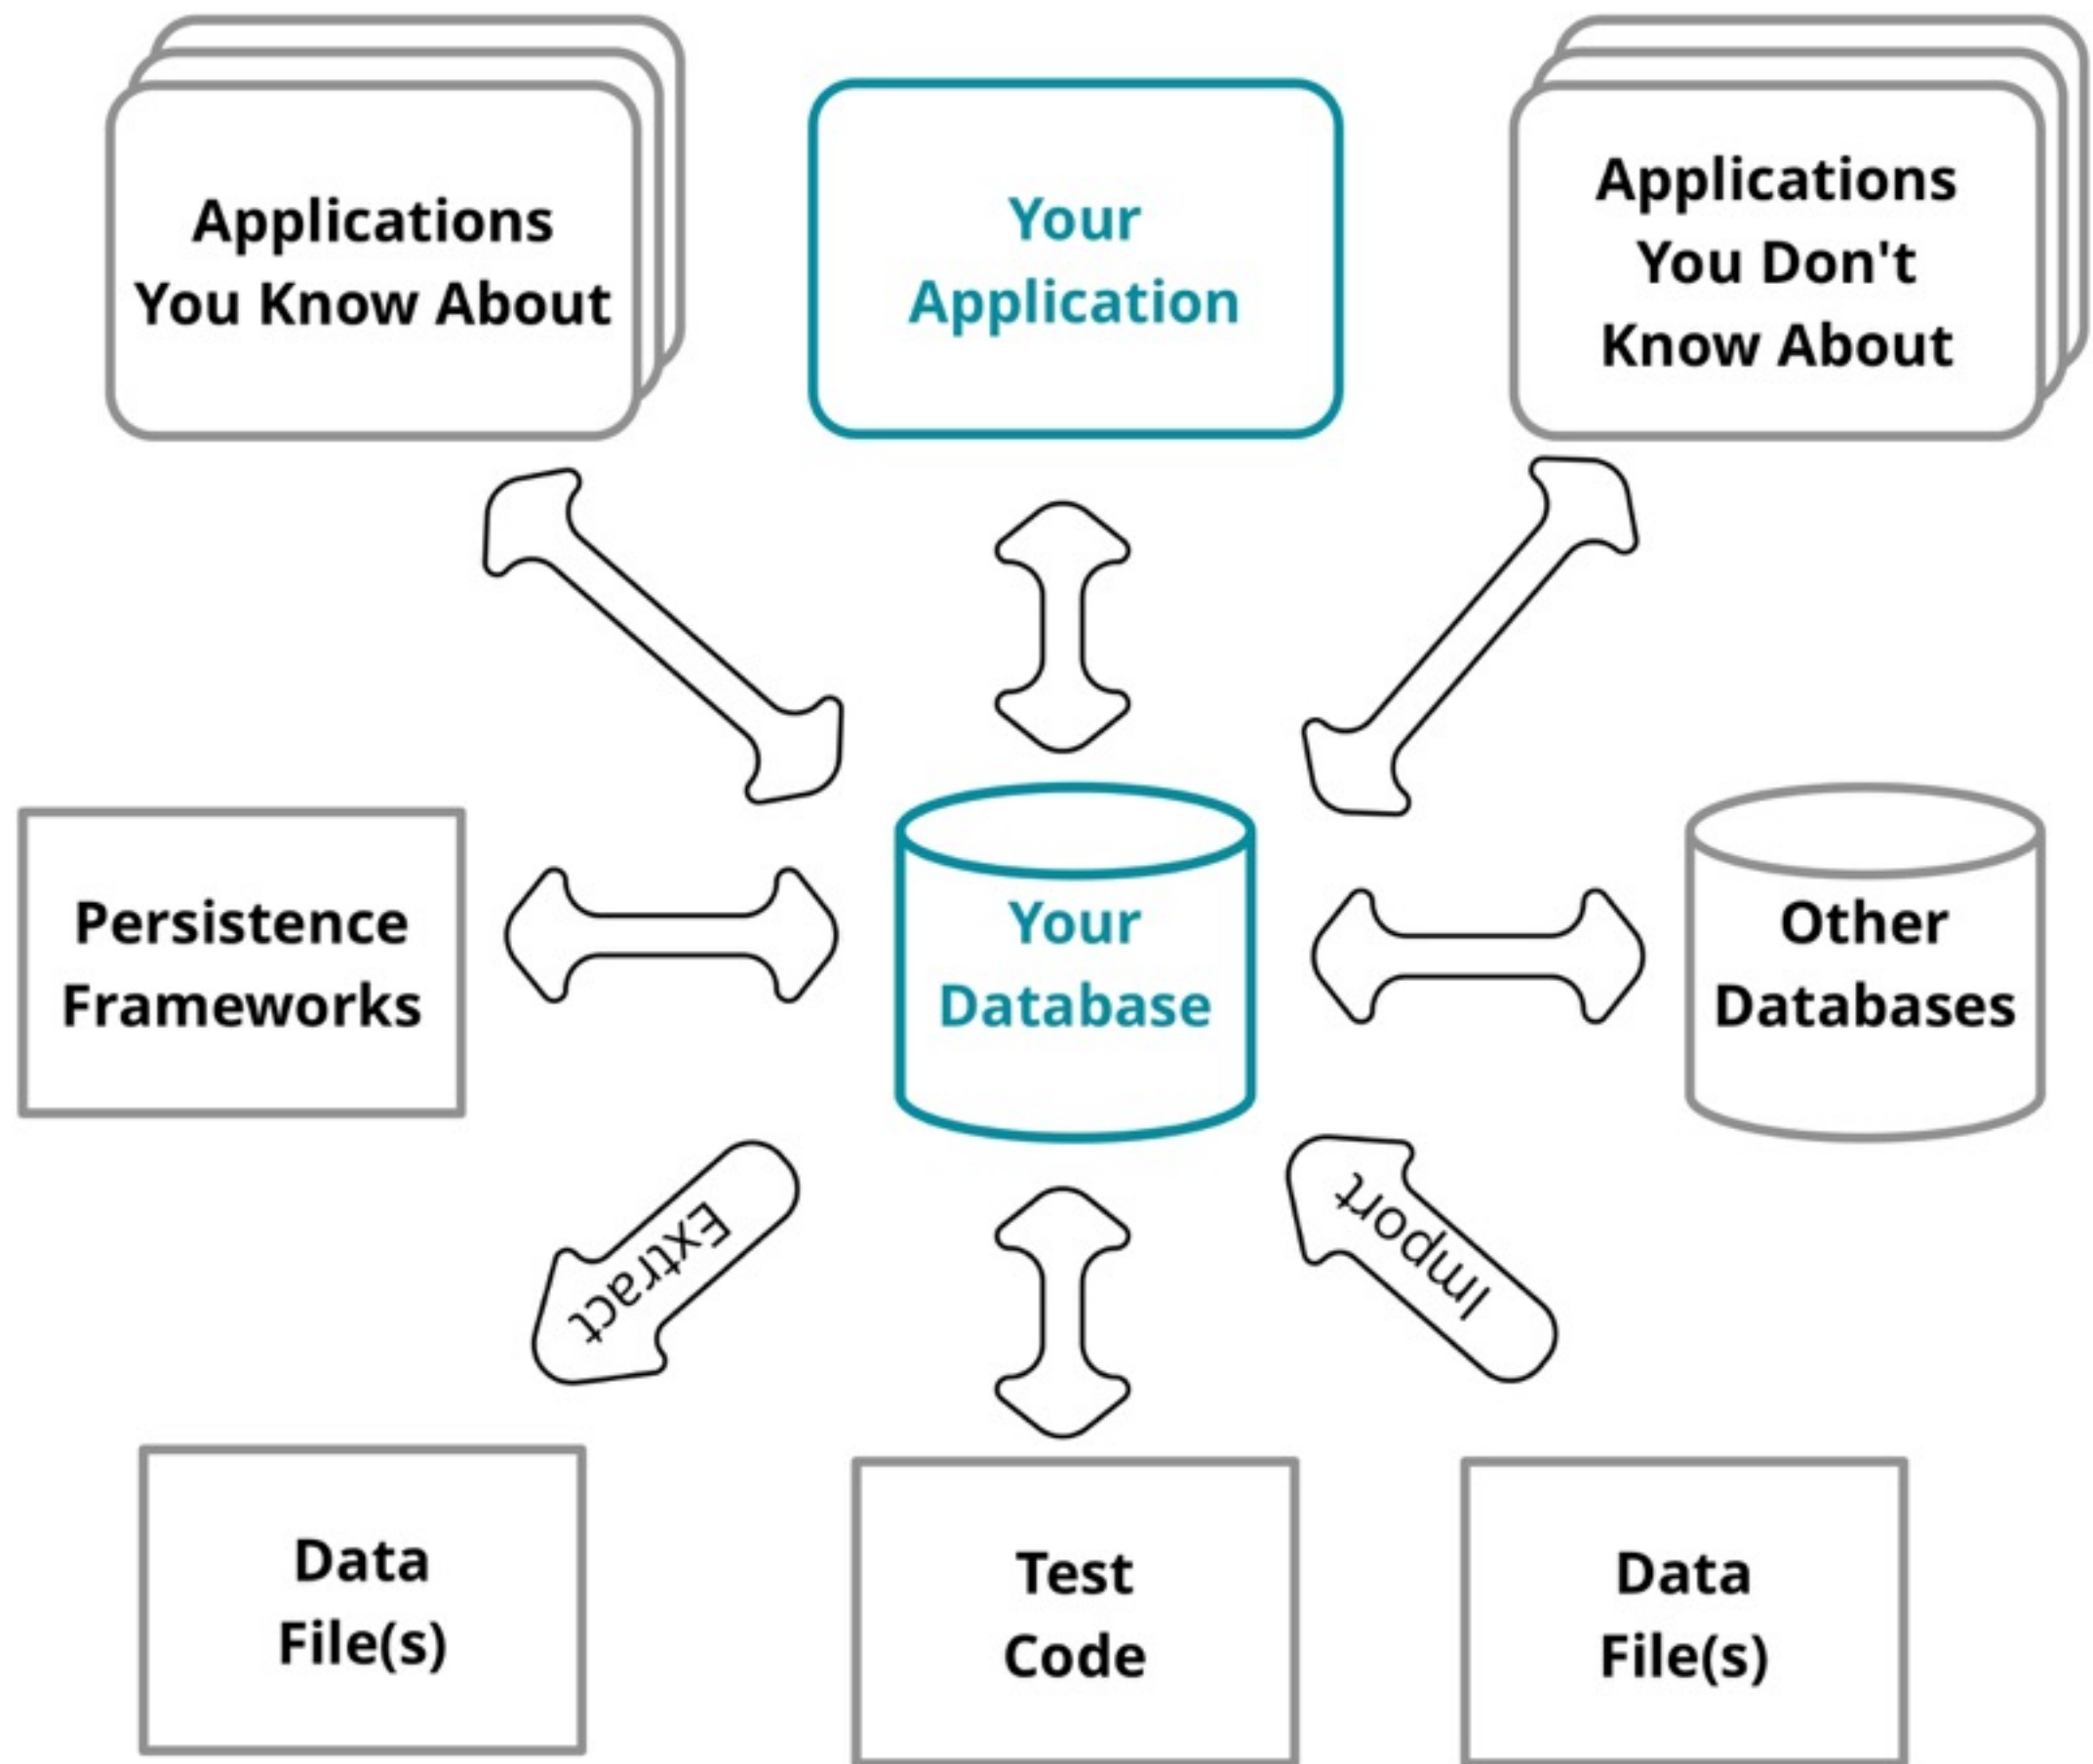
Can your devops process use this to deploy database changes along with code changes?

Database refactoring

<http://databasesrefactoring.com/>

A database refactoring is a small change to your database schema (the DDL, data, and DB code) which improves its design without changing its semantics.

A database refactoring improves its design while retaining both its **behavioral** and **informational semantics**.



Timeline of Change

Deploy new changes, migrate data, put in scaffolding code

Remove old schema, scaffolding code



Implement the refactoring

Transition Period (old and new)

Refactoring completed

Expand

Contract

Keeping Old and New alive

- DB Should be able to handle multiple versions of the application
- Create Interfaces in the database
- Wrap tables with views
- Create calculated columns
- Create triggers to sync data

Expand Contract an example

Customer
CustomerID Name

Starting State

Customer
CustomerID Name FirstName LastName
SynchronizeCustomerName { event = update insert }

Expand State

Customer
CustomerID FirstName LastName

Contracted State

Start

name = "Pramod Sadalage"

Expand

name = "Pramod Sadalage"

firstname = "Pramod"

lastname = "Sadalage"

Contract

firstname = "Pramod"

lastname = "Sadalage"

More on expand contract

Start

name = "Pramod Sadalage"

Expand

Without data migration

name = "Pramod Sadalage"
firstname = null
lastname = null

With data migration

name = "Pramod Sadalage"
firstname = "Pramod"
lastname = "Sadalage"

Contract

firstname = "Pramod"
lastname = "Sadalage"

Simple scenario - DBDeploy

```
ALTER TABLE Customer ADD firstname VARCHAR2(60);  
ALTER TABLE Customer ADD lastname VARCHAR2(60);
```

--//@UNDO

```
ALTER TABLE Customer DROP COLUMN firstname VARCHAR2(60);  
ALTER TABLE Customer DROP COLUMN lastname VARCHAR2(60);
```

With synchronized data

```
ALTER TABLE Customer ADD firstname VARCHAR2(60);
ALTER TABLE Customer ADD lastname VARCHAR2(60);
CREATE OR REPLACE TRIGGER SynchronizeName
BEFORE INSERT OR UPDATE
ON Customer
REFERENCING OLD AS OLD NEW AS NEW
FOR EACH ROW
BEGIN
    IF :NEW.Name IS NULL THEN
        :NEW.Name := :NEW.firstname || ' ' || :NEW.lastname;
    END IF;
    IF :NEW.name IS NOT NULL THEN
        :NEW.firstname := extractfirstname(:NEW.name);
        :NEW.lastname := extractlastname(:NEW.name);
    END IF;
END;
/

--//@UNDO
.....
```

Migrate and Synchronize data

```
ALTER TABLE Customer ADD firstname VARCHAR2(60);  
ALTER TABLE Customer ADD lastname VARCHAR2(60);  
UPDATE Customer set firstname = extractfirstname (name);  
UPDATE Customer set lastname = extractlastname (name);
```

```
CREATE OR REPLACE TRIGGER SynchronizeName  
BEFORE INSERT OR UPDATE
```

```
....
```

```
--//@UNDO
```

```
.....
```

```
UPDATE Customer set name = firstname || ' ' || lastname  
WHERE name IS NULL;  
ALTER TABLE Customer DROP COLUMN firstname;  
ALTER TABLE Customer DROP COLUMN lastname;
```

Contract

ALTER TABLE Customer **SET UNUSED name;**

When
drop takes
forever

—//@UNDO

ALTER TABLE Customer ADD name VARCHAR2(120);
UPDATE Customer set name = firstname || ' ' || lastname
WHERE name IS NULL;

Keep legacy apps happy

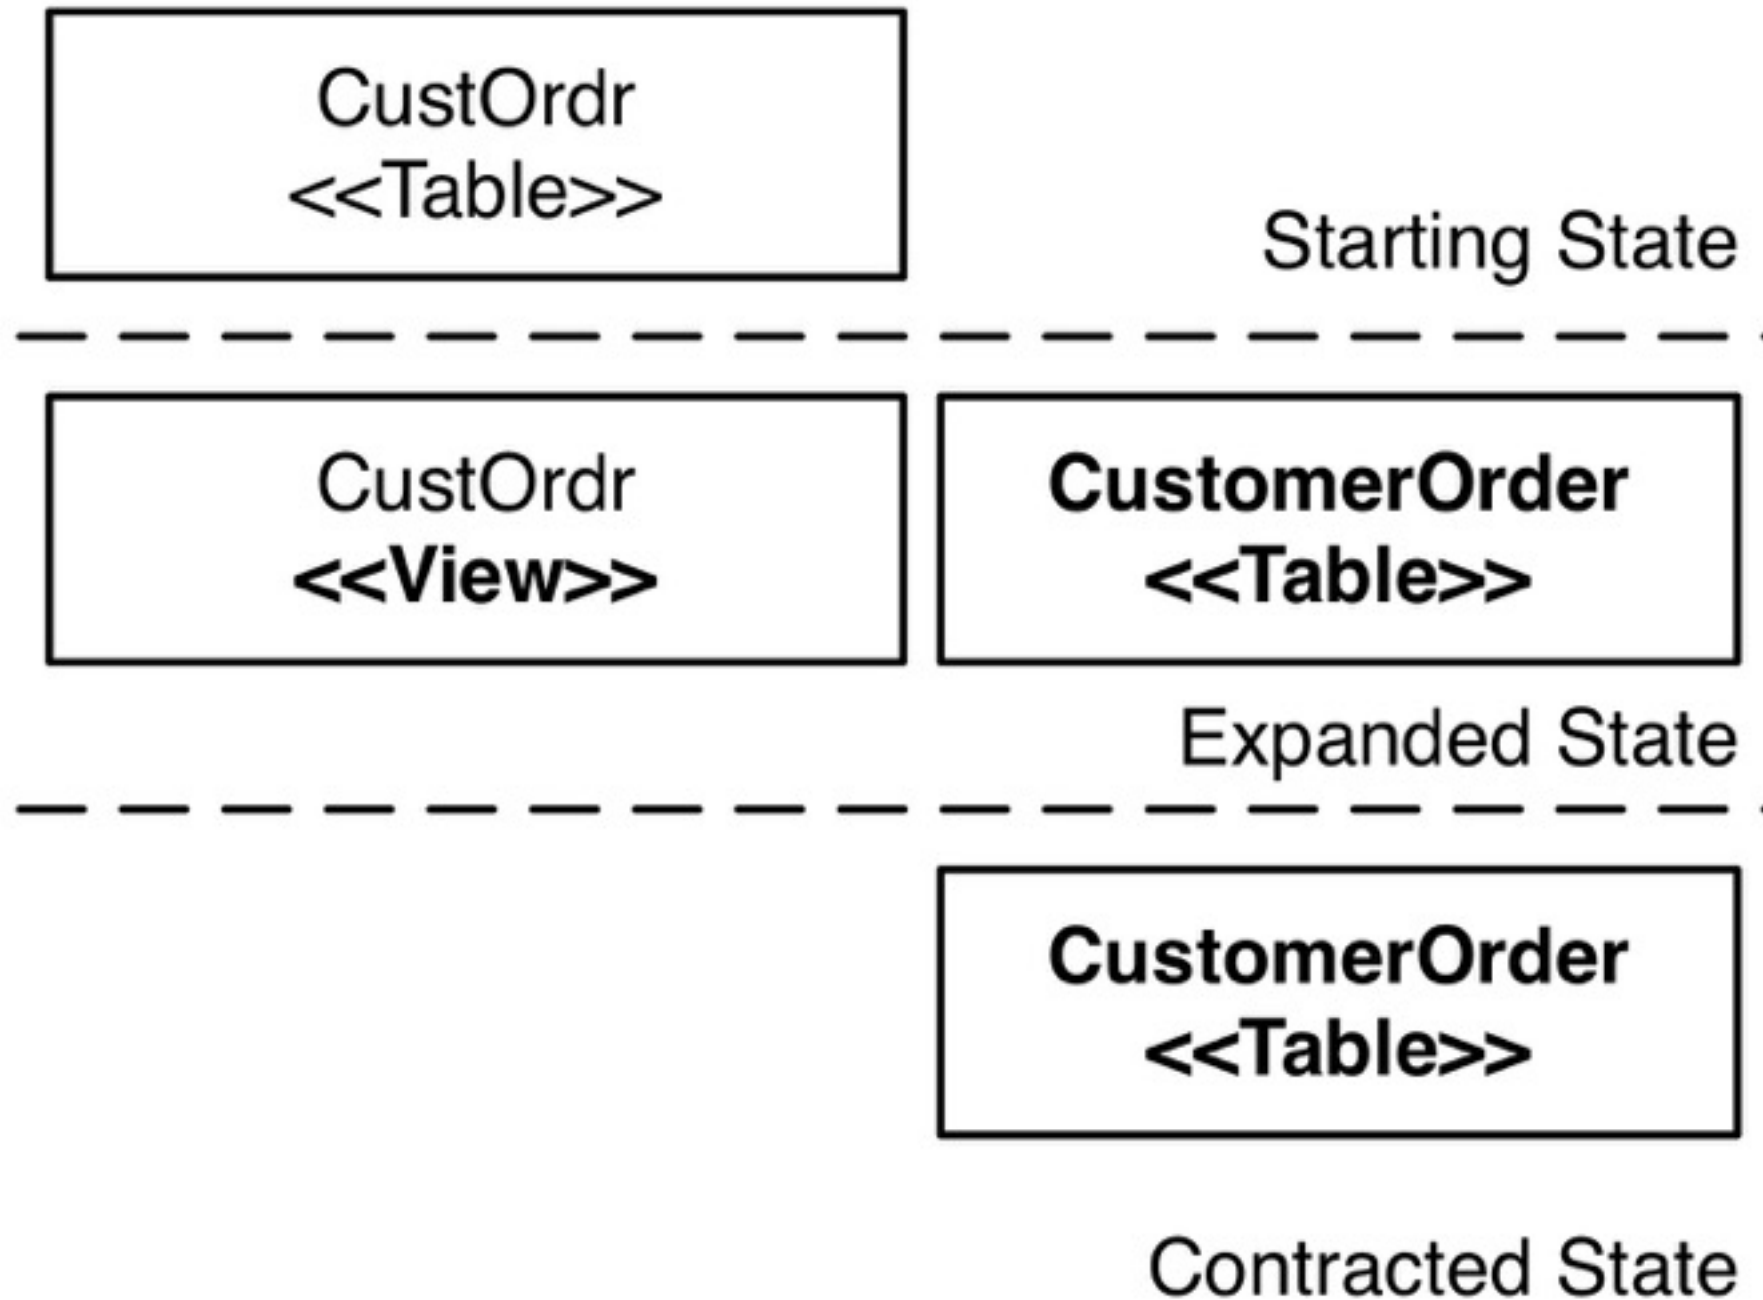
```
ALTER TABLE Customer DROP COLUMN name;  
ALTER TABLE CUSTOMER ADD (name AS  
    (generatename (firstname,lastname))  
    );
```

Virtual column
in Oracle,
Generated
Column in
MySQL.

—//@UNDO

```
ALTER TABEL Customer DROP COLUMN name;  
ALTER TABLE Customer ADD name VARCHAR2(120);  
UPDATE Customer set name = firstname || ' ' || lastname  
WHERE name IS NULL;
```

Another example



Migration script

```
ALTER TABLE custodr rename to customerorder;
```

```
CREATE OR REPLACE VIEW custodr AS  
SELECT custordrid, ponumber, ordrdt, shipdate, sptoadid  
FROM customerorder  
;  
--//@UNDO
```

```
DROP VIEW custodr;  
ALTER TABLE customerorder RENAME TO custodr;
```

data in migrations

```
INSERT INTO businessunit ( businessunitid, name, regionid )
VALUES ( 22, 'John Doe Services', 1 );
DELETE FROM contact ct WHERE NOT EXISTS
(SELECT 1 FROM customer p WHERE ct.contactid=p.contactid)
and EXISTS
(SELECT 1 FROM
    (SELECT customerid, COUNT(*) FROM contact
    WHERE customerid IS NOT NULL GROUP BY customerid
    HAVING COUNT(*) >1) ct2
WHERE ct.customerid=ct2.customerid);
VALUES ( 3, 'Australian Dollar', 'AUD' );
INSERT INTO currency ( currencyid, name, code )
VALUES ( 4, 'EMU Euro', 'EUR' );
```

Shifting to code

Lets refactor our database.

See `migrations/future_migrations`

How would you refactor based on your context?

What are the contexts that you need to think about?



Tips

- Large refactorings are risky
- Sequence of many small refactorings can create the desired change
- Migration scripts should be checked in and run on local dev/ci/qa/uat/prod etc.
- Changes to data are also migrations

Devops for Database

- Devops practices help in Evolving Databases
- Production DBA's are valuable in development
- Automate
- devopsfordba.com

Devops

Developer databases don't have all the
production infrastructure

Partitioning

Linked Databases


Synonyms vs Direct schema access

Continuous Delivery

Deployment should be easy

```
ant -propertyfile qa.properties upgrade
```

```
ant -propertyfile live.properties upgrade
```



Thanks

@pramodsadalage

sadalage.com

databaserefactoring.com

devopsfordba.com